


ACCURACY
How often is the model right overall?



Overall Performance

8 SPAM
5 OK


Predicted

	SPAM	OK
Actual SPAM	TP 5 ✓	3 ✗
Actual OK	2 ✗	TN 3 ✓

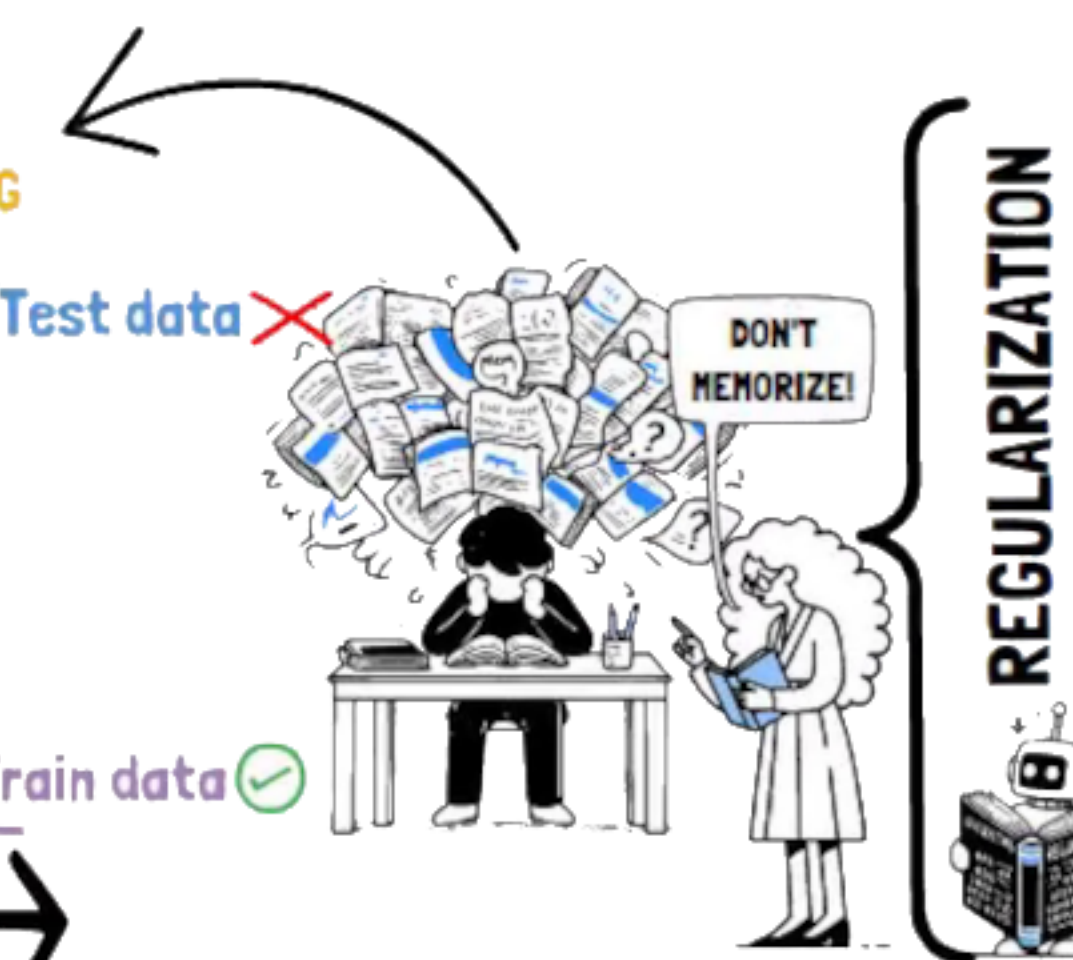




CONFUSION MATRIX



PRECISION
How often is the model right when flagging SPAM?



RECALL
SPAM emails that were caught



L1 / L2
Penalty on complexity of model

DROPOUT
Prevent single neurons from becoming overly dominant